

Japanese, Step by Step, Sign by Sign

A gentle introduction to contemporary standard Japanese,
one sign and one word at a time

Adhithya Rajasekaran

Starter edition — eight short lessons in one cumulative chapter.

Companion practice lessons live alongside this book at
code/learning/human-languages/japanese/lessons/

Copyright © Adhithya Rajasekaran. Licensed under Creative Commons
Attribution-ShareAlike 4.0 International (CC BY-SA 4.0). Japanese is set
in Noto Sans JP (SIL Open Font License).

How to use this starter edition

Begin with no assumed Japanese. Each section takes one small expression, teaches only the writing that expression needs, explains where the word came from, and ends with a short retrieval prompt. There is no kana chart to memorise first.

Japanese uses three writing systems at the same time, and an ordinary sentence contains all three. *Hiragana* carries the grammar, *kanji* carry the content words, and *katakana* marks words borrowed from outside. This edition therefore opens all three in its first chapter rather than finishing one before starting the next, because no real sentence waits for you to.

Two habits are worth forming immediately. First, count beats. Every written kana is one beat of equal length, and length changes meaning: いえ is a house, いいえ is “no”. Second, expect a kanji to have several readings. The character 日 is read *nichi*, *jitsu*, *hi*, *bi*, or *ka* depending on the word it stands in, so this book always teaches the reading a word uses rather than pretending a character has one sound.

Romanization stands beside the Japanese here and will fade in later editions. Where a lesson is teaching beat count it also spells a word out beat by beat, because the macron in *arigatō* hides exactly what is being taught.

A note on what this book does not claim. Japanese shares no ancestor with English, so there is no family of English cousin words to hang new vocabulary on — and none is invented here. What Japanese does have is a real borrowed layer from Chinese, shared with Mandarin and Korean, an internal etymology of its own, and a handful of genuine loans from Portuguese, Dutch, German, and English. Those are used. Nothing else is.

Contents

How to use this starter edition	iii
1 Three Writing Systems in One Doorway	1
1.1 hai	1
1.2 iie	2
1.3 konnichiwa	3
1.4 nihongo	4
1.5 kōhī	6
1.6 arigatō	7
1.7 arigatō gozaimasu	8
1.8 Practice	9
Sources for this starter edition	11

Chapter 1

Three Writing Systems in One Doorway

1.1 はい — yes, in two equal beats

Nothing is assumed yet. Begin with the shortest useful answer in the language, and with the first of its three writing systems.

Script — two signs, two beats

はい — は *ha* + い *i*

Japanese runs left to right. はい is written with two signs of **hiragana**, the everyday script. A hiragana sign is not a letter. It is one whole beat — a **mora** — and normally a consonant and a vowel already fused into one shape. There is no separate *h* to write, because は already is *ha*.

Five vowels, and they stay pure: *a* as in father, *i* as in machine, then *u*, *e*, *o*. So はい is *ha-i*, two beats of the same length, and not the single gliding sound of English “high.”

The word, taken apart

はい means **yes**, and also “here you are,” “I am listening,” and “present.” Its origin is genuinely unsettled: some accounts tie it to 拝 *hai*, “a bow”; others treat it as an old interjection. Neither is proven, so hold it as an open question rather than a fact.

Say plainly what is true here. Japanese shares no ancestor with English, so no cousin word is waiting in your vocabulary to hold はい in place. What this track offers instead is the writing itself, built up one sign at a time, and — later — the real family Japanese does have, through the

characters it borrowed.

Your turn

- *Say it:* はい — *ha-i*, two even beats] [REPEAT x2
- *Read it:* は, then い, left to right
- *Answer:* your name is called → はい

Before you move on

How many beats does はい hold? (Two.) Which sign carries *ha*? (は.)

Source: Wiktionary: はい.

1.2 いいえ — no, and why the extra beat matters

Say はい once and count its two beats out loud. This lesson adds the opposite answer, and one fact about beats that Japanese will never let go of.

Script — one new sign, three beats

いいえ — い *i* + い *i* + え *e*

Only え *e* is new; い you already read in はい. Written out, the word is three signs, so it is three beats: *i-i-e*. The first two are the same sign twice, which is exactly how Japanese writes a long vowel in hiragana — by writing the vowel again.

Beat count is not decoration. いえ (two beats) means **house**. いいえ (three beats) means **no**. Nothing else distinguishes them. Hold the first *i* for two beats and the word lands; clip it and you have said “house.”

The word, taken apart

いいえ is the plain **no** — and, just as often, a modest “**not at all**” used to wave away thanks or a compliment. That second use is the common one in daily speech, which is worth knowing before you meet it as an answer to **thank you** at the end of this chapter.

Your turn

- *Say it:* いいえ — *i-i-e*, three even beats] [REPEAT x2
- *Contrast:* いえ house, いいえ no
- *Choose:* agree → はい; decline → いいえ

Before you move on

Which single sign is new here? (え.) What changes when the long *i* is cut short? (The word becomes いえ, house.)
Source: Wiktionary: いいえ.

1.3 こんにちは — hello, and a sign that lies about its sound

Retrieve はい and the sign は. You are about to meet は again in the daytime greeting — where it is not pronounced *ha*.

Script — four new signs, and one you know

こんにちは — こ *ko* + ん *n* + に *ni* + ち *chi* + は

Four signs are new. こ *ko*, に *ni*, and ち *chi* behave like every hiragana so far: one sign, one consonant-plus-vowel beat.

ん is the exception worth naming. It is the one hiragana that is a bare consonant, and it still takes a full beat of its own. So こんにちは is five signs and five beats: *ko-n-ni-chi-wa*. Give ん its own beat and the greeting sounds Japanese; swallow it and it does not.

Grammar Lens — the sign は, read *wa*

The last sign is は, which you learned as *ha*. Here it is read **wa**. That is not an exception to memorize blindly. は is Japanese's **topic particle** — the little word that marks what a sentence is about, roughly “as for ____.” Japanese spelling was reformed in 1946 to follow pronunciation, but three grammatical particles were left at their older spellings because they are so frequent that respelling them would have rewritten every page in the country. は is one of them. Rule to carry forward: は is *ha* inside a word, and *wa* when it is doing the job of a particle.

Why it's said this way

こんにちは is the daytime greeting — roughly midday to dusk, not morning and not evening. It is a sentence that was abandoned halfway. In full it was 今日は... *konnichi wa...*, “as for today ...”, opening a longer enquiry after someone’s health. The rest fell away and the topic marker stayed, which is why Japan’s most common greeting ends on a grammatical particle instead of a word.

Those two characters, 今日 *konnichi* “this day,” belong to the third writing system, and the next lesson opens it.

Your turn

- *Say it:* こんにちは — *ko-n-ni-chi-wa*, five beats] [REPEAT x2
- *Decide:* は in はい → *ha*; は ending the greeting → *wa*
- *Greet:* someone at midday → こんにちは

Before you move on

Which sign takes a beat without a vowel? (ん.) Why does the greeting end in a particle? (It is the opening of a sentence nobody finishes any more.)

Source: Wiktionary: こんにちは.

1.4 日本語 — three characters, and no single sound each

Say こんにちは. Its history hid two characters, 今日, that hiragana cannot show you. This lesson opens that second system by naming the language itself.

Script — kanji, and the readings problem

日本語 — *ni + hon + go*

These are **kanji**: characters borrowed from Chinese, each carrying a meaning rather than a sound. 日 is sun or day. 本 is origin or root. 語 is language. So 日本 is “sun-origin” — the country the sun rises from — and 日本語 is its language.

Now the fact that has no parallel anywhere else in this curriculum. A kanji does not have *a* sound. 日 is read *nichi*, *jitsu*, *hi*, *bi*, or *ka*, depending on the word it stands in — and in 日本 it is shortened again, to *ni*. There is no sensible way to teach “the sound of 日.” What a lesson can teach is the **word**: you learn 日本語 as *nihongo*, and the reading of each character comes along with it.

Two families of readings explain the spread. The **on** reading is Japan’s version of the Chinese pronunciation, imported with the character. The **kun** reading is a native Japanese word written with a borrowed character that already meant the same thing. 日本語 is on-readings throughout; 今日 in the greeting is irregular even by those standards.

The word, taken apart

Here the cousin web finally works — just not toward English. The on-readings *nichi*, *hon*, *go* are borrowings of Middle Chinese, so 日本語 has genuine relatives: Mandarin writes the same three characters and says *Rìběnyǔ*; Korean borrowed the same compound as *ilbon-eo*. Same characters, related imported sounds, one meaning, three languages. That is a real, checkable connection, and it is the honest replacement for the Latin roots this curriculum leans on elsewhere. English has no share in it.

Your turn

- *Say it*: 日本語 — *ni-ho-n-go*, four beats] [REPEAT x2
- *Name*: 日 sun or day; 本 origin; 語 language
- *Connect*: Mandarin *Rìběnyǔ*, Korean *ilbon-eo*, same three characters

Before you move on

What does 日本 mean literally? (Sun-origin.) Does 日 have one pronunciation? (No — the word it appears in chooses one.)
Source: Wiktionary: 日本語.

1.5 コーヒー — the third system, and a word that really is a cousin

Recall how hiragana writes a long vowel: it writes the vowel twice, as in いいえ. Katakana does it differently, and this word shows how.

Script — katakana, the borrowed-word script

コーヒー — コ *ko* + ィ + ヒ *hi* + ー

This is **katakana**, the third system, and it is a full twin of hiragana: the same beats, the same five vowels, different shapes. Its job is different too. Katakana marks a word as coming from outside — borrowed vocabulary, foreign names, sound effects, scientific terms.

The two scripts are siblings by origin as well as by sound. Hiragana came from whole Chinese characters written in flowing cursive; katakana came from *fragments* of characters, shorthand that monks scribbled beside a text to remember how to read it. ㇿ and ㇿ even trace back to the same character, which is why they still look alike.

ー is the *chōonpu*, the long-vowel bar. It holds the previous vowel for one more beat, and it belongs to katakana almost exclusively. So コーヒー is four beats: *ko-o-hi-i*.

The word, taken apart

コーヒー is **coffee**, and it entered Japanese from Dutch *koffie* through the Nagasaki trade. Dutch had it from Turkish *kahve*, and Turkish from Arabic *qahwa* — which is exactly where English *coffee* comes from as well.

So this word genuinely is a cousin of one you already own. Notice what kind of cousin: not an inherited relative, but the same borrowed object arriving from the same Arabic source by two different roads. That is the honest shape of most Japanese–English overlap. Other passengers on those roads: パン bread, from Portuguese *pão*, and アルバイト a part-time job, from German *Arbeit*.

Your turn

- Say it: コーヒー — *ko-o-hi-i*, four beats] [REPEAT x2

- *Trace it:* Arabic *qahwa* → Turkish *kahve* → Dutch *koffie* → コーヒー
- *Decide:* a borrowed word is written in katakana, not hiragana

Before you move on

Which script signals a borrowed word? (Katakana.) How many beats does コーヒー hold? (Four.)

Source: Wiktionary: コーヒー.

1.6 ありがとう — thanks, from “this hardly ever happens”

Back to hiragana. Five signs are coming, and one small mark that unlocks a whole second row of sounds.

Script — five signs and two small strokes

ありがとう — あ *a* + り *ri* + が *ga* + と *to* + う

あ *a*, り *ri*, と *to* and う *u* are ordinary hiragana. が is the new idea: it is the sign か *ka* with two short strokes added at the top right. That mark is the **dakuten**, and it voices the consonant — the same two strokes turn *ka* into *ga*, *sa* into *za*, *ta* into *da*, *ha* into *ba*. One mark, learned once, doubles how much of the script you can read.

The final う does not sound like *u* here. After と it stretches the *o*, so the word ends *-tō* and holds five beats: *a-ri-ga-to-o*. Hiragana lengthens a vowel by writing another vowel sign, where katakana used the bar ー.

The word, taken apart

Written in kanji, ありがとう is 有り難う: 有る *aru* “to exist” plus 難し *katashi* “difficult.” Its ancestor 有り難し *arigatashi* meant literally “hard to exist” — that is, rare. Rare became precious, precious became grateful, and by the Edo period the word had settled into plain thanks. The *k* of *katashi* softened to *g* inside the compound, which is the same voicing the dakuten writes.

So the memory hook is inside the word, not inside English: saying ありがとう is saying *this was not a thing that had to happen*. No English cousin exists, and inventing one would be worse than admitting that.

Your turn

- *Say it:* ありがとう — *a-ri-ga-to-o*, five beats] [REPEAT x2
- *Build:* か plus the dakuten → が
- *Connect:* 有り難し “hard to exist” → rare → precious → thanks

Before you move on

What did 有り難し mean literally? (Hard to exist.) What does the last う do? (It lengthens the *o*.)
Source: Wiktionary: ありがとう.

1.7 ありがとうございます — politeness you cannot leave out

Say ありがとう once and rebuild が from か. This lesson adds the ending that makes it usable with a stranger, a shopkeeper, or a boss.

Script — five more signs, two of them dakuten

ございます — ご *go* + ざ *za* + い *i* + ま *ma* + す *su*

Two of the five are dakuten forms you can already build: ご is こ with the two strokes, and ざ is さ *sa* with them. い you have read since はい. Only ま *ma* and す *su* are genuinely new.

The whole phrase is ten beats — *a-ri-ga-to-o go-za-i-ma-su* — and the final す is normally said with the *u* barely voiced, close to a plain *s*.

Grammar Lens — politeness is grammar here, not word choice

ございます is a verb. It is the polite form of ある *aru*, “to be” — the same verb that sits inside ありがとう. The route runs ある → ござる → ございます, where the ending -ます is the general politeness marker Japanese attaches to verbs.

That is the difference this chapter has to name. In many languages politeness is a choice between two words for “you.” Japanese instead marks it on the **predicate**, and every sentence has one, so every sentence is marked. There is no neutral setting: ありがとう to a friend and

ありがとうございます to a stranger are not two styles of the same form, they are two different grammatical levels, and choosing wrongly is heard immediately.

At the far end of the system, the verb itself is replaced. 言う *iu* “to say” becomes おっしゃる when the speaker is honouring someone else and 申す when humbling themselves. Same act, three verbs, chosen by who is doing what to whom. That is **keigo**, and this lesson is only its front door.

Your turn

- *Say it:* ありがとうございます — ten beats, final *su* light] [REPEAT x2
- *Choose:* a close friend → ありがとう; a shop assistant → ありがとうございます
- *Build:* こ → ご, さ → ざ

Before you move on

Which verb hides inside ございます? (ある, “to be.”) Can a Japanese sentence avoid marking politeness? (No — the predicate always carries a level.)

Source: Wiktionary: ございます.

1.8 Practice — one daytime exchange, three scripts

Before reading on, retrieve three answers: the daytime greeting, the short yes, and the short no. Nothing new is added here.

The exchange — every line already learned

A: こんにちは。 — *Konnichiwa*. B: こんにちは。 — *Konnichiwa*. A: 日本語? — *Nihongo?* B: はい。 — *Hai*. A: ありがとうございます。 — *Arigatō gozaimasu*. B: いいえ。 — *Iie*.

Three lines of hiragana, one of kanji, and no word that has not been taught. Japanese lets a bare noun with a rising tone stand as a question, so 日本語 with the voice lifted is a complete and natural way to ask “Japanese, then?” — and はい answers it. B’s closing いいえ is the modest “not at all” that waves thanks away.

A single ordinary Japanese sentence mixes the systems: hiragana carries the grammar, kanji carries the content words, katakana carries the borrowings. That is why this chapter opened all three instead of finishing one.

Your turn

- *Run through:* both voices once with the Japanese, once with romanization only
- *Decide:* which line would change if B were a close friend] [PAUSE 3s
- *Swap:* take B’s part while the other lines are read to you

Before you move on

Run the exchange once from memory. If a line stalls, return to that one lesson rather than rereading the chapter.

Sources for this starter edition

Word histories follow Wiktionary's Japanese entries for はい, 今日は, 有難う, and コーヒー. The companion Markdown lessons carry the per-lesson citations and are the canonical short-practice source for this edition; the chapter that follows is generated from them, so the two cannot drift apart.